

tidyverse intro

Jeremy Mikecz

Introduction to Tidyverse

Tidyverse is a system of R packages for data wrangling and analysis. It provides a different method and syntax for working with data tables (often called “data frames” in data science; but, tidyverse dataframes are known as “tibbles”) from base R.

For additional help learning how to use R’s tidyverse system of packages, see:

- Wickham, Çetinkaya-Rundel, and Grolemund, [*R for Data Science*](#)
- Grolemund, [*A Tidyverse Cookbook*](#) (2020).!
- package documentation: [tidyverse](#)
- [tidy cheatsheet](#)
- [*Tidyverse style guide*](#)

For a more detailed introduction to working with tidy data in R see:

- Wickham, Çetinkaya-Rundel, and Grolemund, *R for Data Science*, Ch. 5 “Data Tidying”.
- Grolemund, [*A Tidyverse Cookbook*](#) (2020).
- Silge and Robinson, [*Text Mining with R*](#), Ch. 1 “The tidy text format”.

Getting Started with Tidyverse

Data analysis usually involves working with two-dimensional datasets known as dataframes. There are different ways to work with dataframes in R, these include:

- using core R dataframe functions
- using the **tibble** a type of dataframe used with the tidyverse collection of packages
- using other dataframe types like **data.table** for speed and to facilitate working with large datasets (i.e. with millions of rows)

In these lessons, we will work with `tidyverse` a collection of packages designed for data science. Tidyverse works with `tibbles`, a customized and newer form of dataframes. For more on the differences between tibbles and dataframes see the [explanation here](#). Key packages in tidyverse include (borrowing from Kyle Walker's [Analyzing US Census Data: Methods, Maps, and Models in R](#)):

```
* readr (Wickham and Hester 2021), which contains tools for importing and exporting data  
* dplyr (Wickham et al. 2021), a powerful framework for data wrangling tasks;  
* tidyr (Wickham 2021b), a package for reshaping data;  
* purrr (Henry and Wickham 2020), a comprehensive framework for functional programming a  
* ggplot2 (Wickham 2016), a data visualization package based on the Grammar of Graphics  
## install tidyverse with:  
#install.packages("tidyverse")  
  
## update tidyverse with:  
#tidyverse_update()  
  
## import tidyverse with:  
library(tidyverse)  
  
Warning: package 'tidyverse' was built under R version 4.4.3  
Warning: package 'tibble' was built under R version 4.4.3  
Warning: package 'tidyr' was built under R version 4.4.3  
Warning: package 'readr' was built under R version 4.4.3  
Warning: package 'dplyr' was built under R version 4.4.3  
Warning: package 'stringr' was built under R version 4.4.3  
Warning: package 'forcats' was built under R version 4.4.3  
  
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --  
v dplyr     1.1.4     v readr     2.1.5  
vforcats    1.0.0     v stringr   1.5.1  
v ggplot2   4.0.0     v tibble    3.2.1  
v lubridate 1.9.3     v tidyr    1.3.1  
v purrr    1.0.2  
-- Conflicts ----- tidyverse_conflicts() --  
x dplyr::filter() masks stats::filter()  
x dplyr::lag()   masks stats::lag()  
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to beco
```

```

## see what packages are included with tidyverse:
tidyverse_packages()

[1] "broom"          "conflicted"     "cli"           "dbplyr"
[5] "dplyr"          "dtplyr"         "forcats"       "ggplot2"
[9] "googledrive"   "googlesheets4" "haven"        "hms"
[13] "httr"          "jsonlite"       "lubridate"    "magrittr"
[17] "modelr"         "pillar"         "purrr"        "ragg"
[21] "readr"          "readxl"         "reprex"       "rlang"
[25] "rstudioapi"   "rvest"          "stringr"      "tibble"
[29] "tidyverse"      "xml2"          "tidyverse"

```

Combine functions into a pipe

Using the symbol `|>` we can chain multiple functions together.

Previously, `%>%` was commonly used for pipes. Now, it is recommended to use `|>` instead. In R Studio, you can use the keyboard shortcut CTRL/CMD + SHIFT + M to create a pipe. However, you can update R Studio's default pipe using the instructions in *R for Data Science*, [Ch. 3.4](#).

Each new function in the pipeline operates on the results produced by the previous function.

Chaining functions together in a pipe like this:

```

starwars |>
  group_by(species) |>
  summarise(avg_height = mean(height, na.rm = TRUE)) |>
  arrange(avg_height)

starwars |>
  group_by(species) |>
  summarise(avg_height = mean(height, na.rm = TRUE)) |>
  arrange(avg_height)

# A tibble: 38 x 2
  species      avg_height
  <chr>          <dbl>
1 Yoda's species     66
2 Aleena            79
3 Ewok              88
4 Vulptereen       94
5 Dug               112
6 Xexto             122
7 Droid             131.
8 Toydarian         137

```

```

9 Sullustan      160
10 Toong         163
# i 28 more rows

```

produces the same results as nesting a series of functions within another:

```

arrange(
  summarise(
    group_by(starwars, species),
    avg_height = mean(height, na.rm = TRUE)
  ),
  avg_height
)

```

or calling each function in order:

```

x1 <- starwars
x2 <- group_by(x1, species)
x3 <- summarise(x3, avg_height = mean(height, na.rm = TRUE))
arrange(x3, avg_height)

```

By default `|>` passes the result of the left hand side to the the first unnamed argument of the function on the right hand side. To override this default, use `_` as a placeholder within the function call on the right hand side. `|>` will evaluate `_` as the result of the left hand side, instead of passing the result to the first unnamed argument.

```

starwars |>
  lm(mass ~ height, data = _)

```

```

Call:
lm(formula = mass ~ height, data = starwars)

```

```

Coefficients:
(Intercept)      height
-11.487          0.624

```

```

# the old piping syntax
#starwars %>%
#  lm(mass ~ height, data = .)

```

Tidy Data

From the [Tidyverse cookbook](#):

Data tidying refers to reshaping your data into a tidy data frame or [tibble](#).

Data tidying is an important first step for your analysis because every tidyverse function will expect your data to be stored as **Tidy Data**.

Tidy data is tabular data organized so that:

1. Each column contains a single variable
2. Each row contains a single observation

Tidy data is not an arbitrary requirement of the tidyverse; it is the ideal data format for doing data science with R. Tidy data makes it easy to extract every value of a variable to build a plot or to compute a summary statistic. Tidy data also makes it easy to compute new variables; when your data is tidy, you can rely on R's rowwise operations to maintain the integrity of your observations. Moreover, R can directly manipulate tidy data with R's fast, built-in vectorised observations, which lets your code run as fast as possible.

The definition of Tidy Data isn't complete until you define variable and observation, so let's borrow two definitions from *R for Data Science*:

1. A **variable** is a quantity, quality, or property that you can measure.
2. An **observation** is a set of measurements made under similar conditions (you usually make all of the measurements in an observation at the same time and on the same object).

As you work with data, you will be surprised to realize that what is a variable (or observation) will depend less on the data itself and more on what you are trying to do with it. With enough mental flexibility, you can consider anything to be a variable. However, some variables will be more useful than others for any specific task. In general, if you can formulate your task as an equation (math or code that contains an equals sign), the most useful variables will be the names in the equation.

Create a tibble from scratch

Rarely will you ever create a dataframe or tibble from scratch except to create small practice datasets.

Here, we create a small practice dataset about the popular murder mystery TV show, *Only Murders in the Building*

```
omitb <- tribble(~name, ~occupation, ~apartment, ~is_suspect, ~is_dead,
                  "Charles", "washed-up actor", "14C", "no", "no",
                  "Mabel", "unemployed never-was", "12E", "maybe", "no",
                  "Oliver", "theater director", "10D", "no", "no",
                  "Howard", "childless cat guy", "3D", "yes", "no",
                  "Bunny", "petty despot", "12A", "no", "maybe"
)
omitb
```

```
# A tibble: 5 x 5
  name    occupation      apartment is_suspect is_dead
  <chr>   <chr>          <chr>     <chr>      <chr>
1 Charles washed-up actor    14C       no        no
2 Mabel   unemployed never-was 12E      maybe     no
3 Oliver  theater director    10D       no        no
4 Howard  childless cat guy   3D       yes       no
5 Bunny   petty despot       12A       no        maybe
```

We can add a new row by piping the function `add_row`. To figure out how do so, we can review the [function's documentation](#). Note: the default is to add the new row to the bottom of the tibble. To add it before the first row we need to set the argument `.before = 1`.

```
omitb <- omitb |>
  add_row(name="Jan", occupation = "2nd-chair basoonist",
          is_suspect="yes", is_dead="no", .before=1)
```

```
omitb
```

```
# A tibble: 6 x 5
  name    occupation      apartment is_suspect is_dead
  <chr>   <chr>          <chr>     <chr>      <chr>
1 Jan     2nd-chair basoonist <NA>      yes       no
2 Charles washed-up actor    14C       no        no
3 Mabel   unemployed never-was 12E      maybe     no
4 Oliver  theater director    10D       no        no
5 Howard  childless cat guy   3D       yes       no
6 Bunny   petty despot       12A       no        maybe
```

Above we used the `tribble` function to create a `tibble` (tribble vs. tibble - confusing right?) by using a syntax that lines up the rows neatly in a way that will preview the result.

You can also create the same tibble using the `tibble` function and passing in a list with the values of column:

```
tibble(name = c("Charles", "Mabel", "Oliver", "Howard", "Bunny"),
       occupation = c("washed-up actor", "unemployed never-was", "theater director", "childless cat guy"),
       apartment = c("14C", "12E", "10D", "3D", "12A"),
       is_suspect = c("no", "no", "no", "yes", "no"),
       is_dead = c("no", "no", "no", "no", "maybe"),
       )
```



```
# A tibble: 5 x 5
  name    occupation      apartment is_suspect is_dead
  <chr>   <chr>          <chr>     <chr>      <chr>
1 Charles washed-up actor    14C       no        no
2 Mabel   unemployed never-was 12E       no        no
```

```

3 Oliver theater director      10D      no      no
4 Howard childless cat guy    3D      yes      no
5 Bunny petty despot         12A      no      maybe

```

Dataframes vs. tibbles

```
df <- read.csv("../data/census1970.csv")
```

Convert a dataframe into a tibble

```
head(df)
```

	rownum	STATE	COUNTY		NAME	TOTPOP	WPOP	NEGTOT	OTHRACES	MTOT				
1	1	1	10	FAIRFIELD		792814	732304	56408	4102	381603				
2	2	1	30	HARTFORD		816737	758086	54645	4006	394728				
3	3	1	50	LITCHFIELD		144091	142649	1126	316	70176				
4	4	1	70	MIDDLESEX		114816	111102	3327	387	56465				
5	5	1	90	NEW HAVEN		744948	684743	56630	3575	359204				
6	6	1	110	NEW LONDON		230348	221073	7390	1885	115752				
		M04	M56	M79	M1013	M14	M15	M1617	M1819	M20	M21	M2224	M2534	M3544
1	32152	15536	24082	33345	8129	8059	15054	11245	5006	4485	13869	44111	48524	
2	34432	15927	24454	33771	8118	7984	15238	11862	5168	5129	17665	50392	47027	
3	6172	2833	4302	5993	1491	1488	2784	1975	753	719	2631	8513	7935	
4	4952	2305	3506	4709	1133	1116	2017	1883	825	840	2404	7483	6622	
5	31329	14023	21708	29890	7122	7156	13528	12637	5528	5240	16988	44185	40199	
6	10946	4976	7329	9386	2192	2186	3964	4027	2274	2470	7584	16524	12821	
		M4554	M5559	M6061	M6264	M6574	M75	FTOT	F04	F56	F79	F1013	F14	F15
1	51181	21292	7003	9010	18823	10697	411211	30932	15026	23313	32321	7773	7446	
2	50696	20465	6991	8801	19346	11262	422009	33550	15350	23335	32498	8082	7688	
3	8618	3960	1398	1819	4289	2503	73915	5896	2778	4196	5636	1423	1403	
4	6733	2825	1036	1304	3033	1739	58351	4845	2209	3362	4553	1029	1083	
5	45636	18725	6420	8196	19212	11482	385744	30522	13707	20665	28656	7213	6771	
6	12164	4783	1828	2179	5119	3000	114596	10492	4623	7098	8922	2142	2010	
		F1617	F1819	F20	F21	F2224	F2534	F3544	F4554	F5559	F6061	F6264	F6574	F75
1	14651	11814	5736	5411	16418	49023	51666	54853	22422	7681	10120	26181	18424	
2	14579	13353	6771	6786	20770	51528	49512	53062	21459	7759	10232	27049	18646	
3	2602	1908	884	873	3038	8688	8098	9366	4243	1536	2080	5358	3909	
4	1953	1508	785	799	2741	7527	6593	7239	2887	1104	1408	3829	2897	
5	13060	11985	6097	6060	18958	46105	43485	49626	20698	7375	9759	26603	18399	
6	3880	3890	2047	2043	5826	14816	12645	12821	5094	1939	2479	6794	5035	
		NEGMTOT	NEGM04	NEGM514	NEGM1524	NEGM2534	NEGM3544	NEGM4554	NEGM5564	NEGM65				
1	26161	3344	6917		4332		3977		2943		2262		1427	959
2	26069	3507	7105		4427		3790		2845		2281		1248	866
3	592	59	137		97		69		76		57		52	45

4	1630	203	367	358	271	169	123	71	68	
5	26538	3665	7537	4888	3434	2745	2140	1188	941	
6	3644	416	979	734	465	462	310	160	118	
	NEGFTOT	NEGF04	NEGF514	NEGF1524	NEGF2534	NEGF3544	NEGF4554	NEGF5564	NEGF65	
1	30247	3513	7143	5187	4901	3450	2783	1848	1422	
2	28576	3442	6717	5582	4574	3227	2385	1380	1269	
3	534	62	98	87	62	71	43	68	43	
4	1697	202	433	369	262	167	117	72	75	
5	30092	3917	7436	5544	4565	3386	2498	1460	1286	
6	3746	431	939	744	510	454	348	158	162	
	TOTPOP2	HHPOP	HHHEADS	HHPRIM	HHMHEAD	HHFHEAD	HHFSPOUS	HHOTHREL	HHUNREL	GROUP
1	792814	778212	243806	41569	181747	20490	176172	346686	11548	14602
2	816737	799712	255437	48242	184924	22271	179067	352227	12981	17025
3	144091	142582	45550	8085	34330	3135	33245	62151	1636	1509
4	114816	109867	34758	5830	26409	2519	25706	48076	1327	4949
5	744948	725581	231754	42156	168111	21487	162404	321253	10170	19367
6	230348	214257	67618	11833	49180	6605	47778	96592	2269	16091
	INMATES	GROUPOTH	NEGTOT2	HHNEG	HHNEGHEA	HHNEGPRI	HHNEGMHE	HHNEGFHE	HHNEGFSP	
1	6556	8046	56408	55207	15685	3162	8879	3644	8260	
2	8116	8909	54645	53442	15533	3408	8322	3803	7659	
3	878	631	1126	1075	306	62	214	30	193	
4	3079	1870	3327	2952	819	178	480	161	453	
5	8467	10900	56630	55416	15613	3096	8481	4036	7948	
6	3736	12355	7390	6793	1857	329	1128	400	1042	
	HHNEGOTR	HHNEGUNR	GROUPNEG	INMATNEG	GROUNEGO	REGION1	REGION2	STATEFIP	FIPS	
1	29100	2162	1201	582	619	1	1	9	9001	
2	28526	1724	1203	656	547	1	1	9	9003	
3	512	64	51	29	22	1	1	9	9005	
4	1592	88	375	276	99	1	1	9	9007	
5	30421	1434	1214	742	472	1	1	9	9009	
6	3794	100	597	298	299	1	1	9	9011	
	LEVEL									
1	1									
2	1									
3	1									
4	1									
5	1									
6	1									

```
df2tib <- as_tibble(df)
```

3. Tidyverse / dplyr verbs for Data Wrangling

To modify and “wrangle” datasets, the tidyverse commonly uses the following verbs / functions:

- **select()**: to select specific columns by their names or data types
- **arrange()**: to order rows by one or more columns
- **rename()**: to rename columns
- **mutate()**: to create columns
- **filter()**: to filter out rows by a given condition
- **distinct()**: to keep only distinct / unique rows
- **gather()**: to make “wide” data longer
- **spread()**: to make “long” data wider
- **separate()**: to split a single column into multiple columns
- **unite()**: to combine multiple columns into one

[More tidyverse verbs and functions here.](#)

Common Data Science Operations in Tidyverse

Task	Python - Pandas	Base R	Tidyverse (R)
Sort table by column	<code>df.sort_values('col')</code>	<code>sort(df\$col)</code>	<code>df %>% arrange(col)</code>
Filter rows	<code>df.loc[df['col'] > 0,:]</code>	<code>df[df[col] > 0,</code> <code>]</code>	<code>df %>% filter(col > 0)</code>
Subset columns			<code>df %>% select(col)</code>
Create a new column	<code>df.loc[:, "square"] = df[col] ** 2</code>	<code>df\$square = df[col] ^2</code>	<code>df %>% mutate(square = col ^ 2)</code>
Perform calculation on column	<code>sum(df["col"])</code>	<code>sum(df\$col)</code>	<code>df %>% summarise(sumcol = sum(col))</code>
Delete duplicate rows	<code>df.drop_duplicates()</code>	<code>df[!duplicated(df)]</code> <code>]</code>	<code>df %>% distinct()</code>
Group/Apply/Combine	<code>df.groupby("col").sum()</code>	<code>by(df, col,</code> <code>sum(x))</code> <code>???</code>	<code>df %>% group_by(col) %>% summarise(groupsum = sum(col2))</code>
Reshape data from wide to long	<code>pd.wide_to_long()</code>	<code>melt()</code>	<code>use pivot_longer()</code>

Sort Table (tibble)

```
starwars |>
  arrange(homeworld) #descending order: arrange(desc(homeworld))
```

```

# A tibble: 87 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex   gender
  <chr>     <int> <dbl> <chr>       <chr>       <chr>       <dbl> <chr> <chr>
1 Leia Or~    150    49 brown     light      brown        19 female feminine
2 Bail Pr~    191    NA black     tan        brown        67 male   masculine
3 Raymus ~   188    79 brown     light      brown        NA male   masculine
4 Ratts T~    79     15 none     grey, blue unknown    NA male   masculine
5 Lobot      175    79 none     light      blue         37 male   masculine
6 Jek Ton~   180    110 brown    fair       blue         NA <NA> <NA>
7 Nute Gu~   191    90 none     mottled   grey red     NA male   masculine
8 Ki-Adi~~  198     82 white    pale       yellow     92 male   masculine
9 Mas Ame~   196    NA none     blue       blue        NA male   masculine
10 Mon Mot~  150    NA auburn   fair       blue        48 female feminine
# i 77 more rows
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>

```

Filter Table

```

starwars |>
  filter(species=="Droid")

# A tibble: 6 x 14
  name      height  mass hair_color skin_color eye_color birth_year sex   gender
  <chr>     <int> <dbl> <chr>       <chr>       <chr>       <dbl> <chr> <chr>
1 C-3PO     167    75 <NA>      gold       yellow        112 none   masculine
2 R2-D2     96     32 <NA>      white, blue red        33 none   masculine
3 R5-D4     97     32 <NA>      white, red red        NA none   masculine
4 IG-88    200    140 none     metal       red         15 none   masculine
5 R4-P17    96     NA none     silver, red red, blue    NA none   feminine
6 BB8      NA     NA none     none       black        NA none   masculine
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>

```

Subset Tibble by Column Names

```

starwars |>
  select(name, homeworld, species) #reorder names here to reorder columns

# A tibble: 87 x 3
  name           homeworld species
  <chr>          <chr>     <chr>
1 Luke Skywalker Tatooine Human
2 C-3PO          Tatooine Droid

```

```

3 R2-D2           Naboo   Droid
4 Darth Vader    Tatooine Human
5 Leia Organa    Alderaan Human
6 Owen Lars      Tatooine Human
7 Beru Whitesun Lars Tatooine Human
8 R5-D4           Tatooine Droid
9 Biggs Darklighter Tatooine Human
10 Obi-Wan Kenobi Stewjon  Human
# i 77 more rows

```

For a range of columns

```

starwars |>
  select(name:hair_color, species)

# A tibble: 87 x 5
  name          height  mass hair_color species
  <chr>        <int> <dbl> <chr>    <chr>
1 Luke Skywalker 172     77  blond    Human
2 C-3PO          167     75 <NA>     Droid
3 R2-D2          96      32 <NA>     Droid
4 Darth Vader    202     136 none    Human
5 Leia Organa    150     49 brown   Human
6 Owen Lars      178     120 brown, grey Human
7 Beru Whitesun Lars 165     75 brown   Human
8 R5-D4          97      32 <NA>     Droid
9 Biggs Darklighter 183     84 black   Human
10 Obi-Wan Kenobi 182     77 auburn, white Human
# i 77 more rows

```

Rename columns

```

starwars |>
  rename(character = name, planet = homeworld)

# A tibble: 87 x 14
  character      height  mass hair_color skin_color eye_color birth_year sex
  <chr>        <int> <dbl> <chr>    <chr>    <chr>        <dbl> <chr>
1 Luke Skywalker 172     77  blond    fair      blue         19   male
2 C-3PO          167     75 <NA>     gold     yellow       112  none
3 R2-D2          96      32 <NA>     white, bl~ red        33  none
4 Darth Vader    202     136 none    white     yellow      41.9 male
5 Leia Organa    150     49 brown   light     brown       19  fema-
6 Owen Lars      178     120 brown, gr~ light     blue        52   male
7 Beru Whitesun ~ 165     75 brown   light     blue        47  fema-

```

```

8 R5-D4          97    32 <NA>      white, red red      NA  none
9 Biggs Darkligh~ 183    84 black      light      brown      24  male
10 Obi-Wan Kenobi 182    77 auburn, w~ fair      blue-gray     57  male
# i 77 more rows
# i 6 more variables: gender <chr>, planet <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>
starwars$vehicles[1]

[[1]]
[1] "Snowspeeder"           "Imperial Speeder Bike"

```

Split-Apply-Combine

```

starwars |>
  group_by(species)

# A tibble: 87 x 14
# Groups:   species [38]
  name    height  mass hair_color skin_color eye_color birth_year sex   gender
  <chr>    <int> <dbl> <chr>       <chr>       <chr>       <dbl> <chr> <chr>
  1 Luke Sk~    172    77 blond      fair        blue        19  male  masculin~
  2 C-3PO       167    75 <NA>       gold        yellow      112  none  masculin~
  3 R2-D2        96    32 <NA>       white, bl~ red        33  none  masculin~
  4 Darth V~     202   136 none      white        yellow      41.9 male  masculin~
  5 Leia Or~     150    49 brown      light        brown      19  fema~ feminin~
  6 Owen La~     178   120 brown, gr~ light        blue       52  male  masculin~
  7 Beru Wh~     165    75 brown      light        blue       47  fema~ feminin~
  8 R5-D4        97    32 <NA>       white, red red      NA  none  masculin~
  9 Biggs D~     183    84 black      light      brown      24  male  masculin~
  10 Obi-Wan~    182    77 auburn, w~ fair      blue-gray     57  male  masculin~
# i 77 more rows
# i 5 more variables: homeworld <chr>, species <chr>, films <list>,
#   vehicles <list>, starships <list>
starwars |>
  group_by(species) |>
  summarise(avg_height = mean(height, na.rm = TRUE))

# A tibble: 38 x 2
  species  avg_height
  <chr>        <dbl>
  1 Aleena      79
  2 Besalisk    198
  3 Cerean      198
  4 Chagrian    196

```

```

5 Clawdite      168
6 Droid         131.
7 Dug           112
8 Ewok          88
9 Geonosian    183
10 Gungan       209.
# i 28 more rows

starwars |>
  group_by(species) |>
  summarise(avg_height = mean(height, na.rm = TRUE)) |>
  arrange(avg_height)

# A tibble: 38 x 2
  species      avg_height
  <chr>        <dbl>
1 Yoda's species     66
2 Aleena            79
3 Ewok              88
4 Vulptereen       94
5 Dug               112
6 Xexto             122
7 Droid             131.
8 Toydarian         137
9 Sullustan         160
10 Toong            163
# i 28 more rows

```

💡 Exercise

In a new Quarto document, choose another built-in dataset (run `data()` to see what datasets are available).

1. examine the dataset using the functions demonstrated in the previous notebook (Week 1 / Notebook 01).
2. Pose a question you would like to answer about this dataset?
3. To answer this question re-shape, filter, sort, and aggregate the dataset using the tidyverse functions demonstrated above.
4. Write a brief paragraph explaining what you learned from this preliminary analysis.
5. **To complete this assignment: submit a pdf of this document. You can**